

SDS 383D: Exercises 3 – Linear smoothing and Gaussian processes

February 23, 2017

Professor Scott

Spencer Woody

Problem 1

Basic Concepts

(A)

Problem 2

Curve fitting by linear smoothing

In this problem, consider a general nonlinear regression with one predictor and one response, $y_i = f(x_i) + \epsilon_i$, where ϵ_i are mean-zero random variables.

- (A) For now, consider a linear regression on a response y_i with one predictor x_i , and both y_i and x_i have had their means subtracted, so the $y_i = \beta x_i + \epsilon_i$. Define $S_x := \sum_{i=1}^n x_i^2$. The least squares estimate for the coefficient, from Exercises 1, is

$$\begin{aligned}\hat{\beta} &= (X^T X)^{-1} X^T y \\ &= (x^T x)^{-1} x^T y \\ &= \frac{\sum_{i=1}^n x_i \cdot y_i}{\sum_{i=1}^n x_i^2} \\ &= \frac{\sum_{i=1}^n x_i \cdot y_i}{S_x} \\ &= \sum_{i=1}^n \frac{x_i}{S_x} \cdot y_i.\end{aligned}$$

So now our prediction $y^*|x^*$ is,

$$\begin{aligned}\hat{y}^* &= \hat{f}(x^*) \\ &= \hat{\beta} x^* \\ &= \left(\sum_{i=1}^n \frac{x_i}{S_x} \cdot y_i \right) \cdot x^* \\ &= \sum_{i=1}^n \left(\frac{x_i}{S_x} \cdot x^* \right) \cdot y_i,\end{aligned}$$

which we recognize as being in the form of the general *linear smoother*

$$\hat{f}(x^*) = \sum_{i=1}^n w(x_i, x^*) \cdot y_i$$

for some weight function $w(x_i, x^*)$. In particular, the weight function for linear regression gives weight to each y_i proportional to the value of x_i . Contrast this with the k -nearest neighbors smoothing weight function,

$$w_K(x_i, x^*) = \begin{cases} 1/K & \text{if } x_i \text{ is one of the } K \text{ closest sample points to } x^*, \\ 0 & \text{otherwise} \end{cases},$$

which gives *equal* weight to y_i s but *only* to the k -nearest neighbors of x^* .

- (B) Now we have the very general weight function

$$w(x_i, x^*) = \frac{1}{h} \cdot K\left(\frac{x_i - x^*}{h}\right)$$

where $K(\bullet)$ is some kernel function. The script `myfuns03.R` in the appendix shows an R function for linear smoothing, as well functions for the uniform and Gaussian kernels. Figure 1 shows an example of smoothing with a bandwidth of 0.75 for a cubic function $f(x)$ with iid residuals from the $\mathcal{N}(0, 15^2)$ distribution.

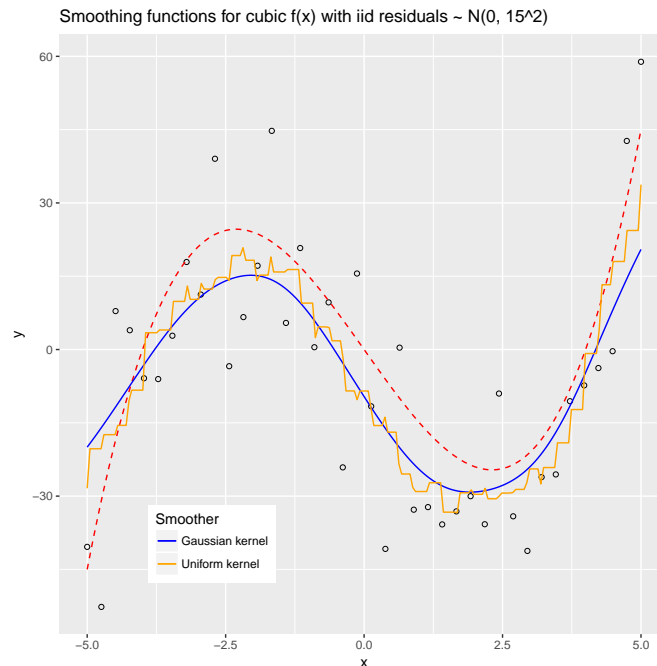


Figure 1: Uniform and Gaussian kernel smoothing for $y = f(x) + e$, $f(x) = x(x - 4)(x + 4)$, $h = 0.75$

Problem 3

Cross validation

- (A) See attached R code for a script to return prediction error estimates for smoothing given a specified choice of bandwidth, h .
- (B) For this exercise, I produced 500 data points on the x -space $[0, 1]$ from a sine function $f(x)$ with a given period and set the amplitude, and added Gaussian noise with a given standard deviation. Then I used 5-fold cross validation to select the optimal bandwidth for that given period and standard deviation of noise term. Figure 2 shows the optimal bandwidths for period ranging from 0.1 to 1, and standard deviation ranging from 0.001 to 0.5, and Figure 3 shows four example. The highest bandwidths are chosen for functions with high “wigglyness” and high noise, and the smallest bandwidths are chosen for functions with low “wigglyness” and low noise. This makes sense. As the frequency increases (i.e., period decreases) then we need a tighter bandwidth because the value of the function is fluctuating at a greater rate. As noise increases, we need a greater bandwidth to smooth out the noise. Furthermore, we can see that in all cases we recover the underlying function pretty well.
- (C) I’ll get around to this eventually....

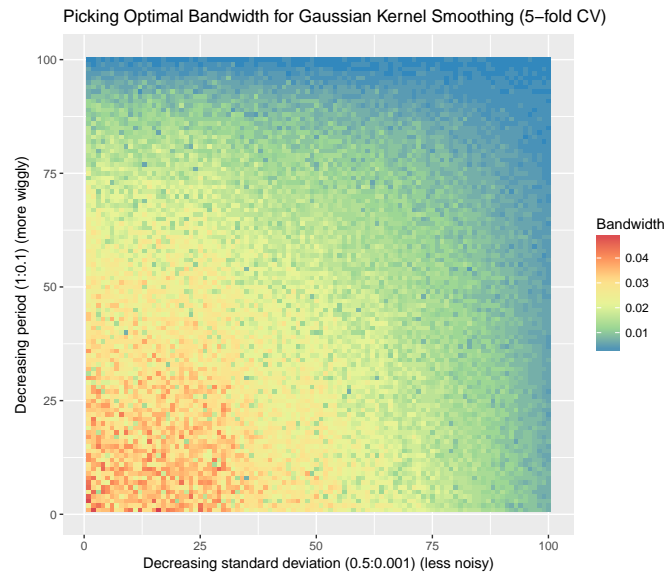


Figure 2: Optimal bandwidths for varying periods and standard deviations

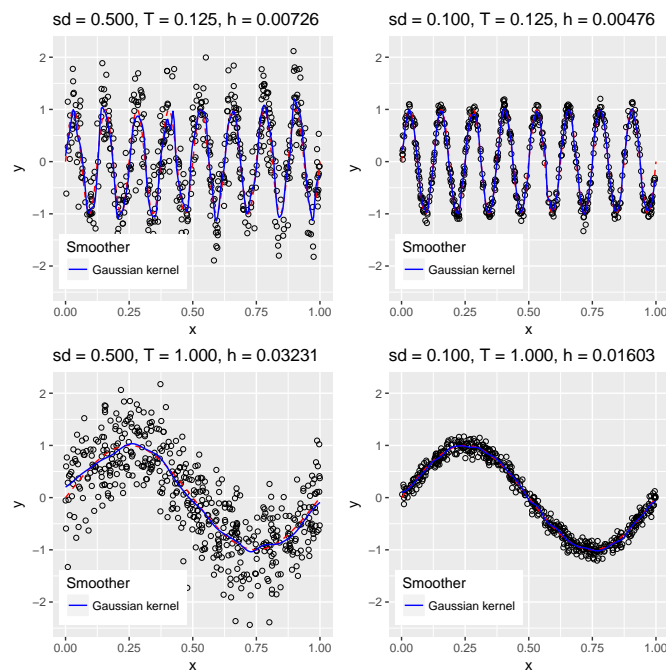


Figure 3: 2×2 example with fitted curves

Problem 4

Local polynomial regression

(A) Define

$$\begin{aligned} g_x(u; a) &= a_0 + \sum_{k=1}^D a_k (u - x)^k \\ &= \begin{cases} \sum_{j=0}^{D+1} a_j (u - x)^j & \text{if } u \neq x \\ a_0 & \text{if } u = x \end{cases}. \end{aligned}$$

The coefficients of a will come from the weighted least squares problem

$$\hat{a} = \arg \min_{a \in \mathcal{R}^{D+1}} \sum_{i=1}^n w_i [y_i - g_x(x_i, a)]^2 \quad (1)$$

Furthermore, define the matrix R_x whose (i, j) element is $(x_i - x)^2$. Then the estimate $\hat{f}(x)$ will be $R_x \hat{a}$. The solution of \hat{a} may be found with

$$\begin{aligned} \hat{a} &= \arg \min_{a \in \mathcal{R}^{D+1}} (y - R_x a)^T W (y - R_x a) \\ &= (R_x^T W R_x)^{-1} R_x^T W y \end{aligned}$$

where $W = \text{diag}(w_1, \dots, w_n)$ following the same argument to find the WLS estimate of linear regression from Exercises 1.

(B)

(C) The estimate of f at x , up to a proportionality is

$$\begin{aligned} \hat{f}(x) &\propto \sum_{i=1}^n w_i g_x(x_i; \hat{a}) \\ &= \mathbf{1}^T W R_x \hat{a} \\ &= \mathbf{1}^T W R_x (R_x^T W R_x)^{-1} R_x^T W y \\ &= b_x^T y \end{aligned}$$

if we define the vector $b_x^T = \mathbf{1}^T W R_x (R_x^T W R_x)^{-1} R_x^T W$. Note that we need to normalize our estimate so that the “weights” sum to one, so now our estimate is exactly

$$\hat{f}(x) = \frac{b_x^T y}{\|b_x\|_1}.$$

(D) With H a smoothing matrix (or “hat matrix”), let $r = y - Hy$ be the vector of residuals. If the random vector x with mean vector μ and covariance matrix Σ , then $E(x^T Q x) = \text{tr}(Q \Sigma) + \mu^T Q \mu$. By assumption,

$E(\mathbf{y}) = f(\mathbf{x})$ and $\text{cov}(\mathbf{y}) = \sigma^2 I$ Then,

$$\begin{aligned}
 E\left(\|r^2\|_2^2\right) &= E\left((\mathbf{y} - H\mathbf{y})^T(\mathbf{y} - H\mathbf{y})\right) \\
 &= E\left(\mathbf{y}^T \mathbf{y} - 2\mathbf{y}^T H\mathbf{y} + \mathbf{y}^T H^T H\mathbf{y}\right) \\
 &= E\left(\mathbf{y}^T \mathbf{y}\right) - 2E\left(\mathbf{y}^T H\mathbf{y}\right) + E\left(\mathbf{y}^T H^T H\mathbf{y}\right) \\
 &= \left(\text{tr}\left[I\sigma^2 I\right] + f(\mathbf{x})^T f(\mathbf{x})\right) - 2\left(\text{tr}\left[H^T \sigma^2 I\right] + f(\mathbf{x})^T H^T f(\mathbf{x})\right) + \left(\text{tr}\left[H^T H\sigma^2 I\right] + f(\mathbf{x})^T H^T H f(\mathbf{x})\right) \\
 &= \left(n\sigma^2 + f(\mathbf{x})^T f(\mathbf{x})\right) - 2\left(\sigma^2 \text{tr}[H] + f(\mathbf{x})^T H^T f(\mathbf{x})\right) + \left(\sigma^2 \text{tr}[H^T H] + f(\mathbf{x})^T H^T H f(\mathbf{x})\right) \\
 &= \left(n - \text{tr}[H] + \text{tr}[H^T H]\right) \sigma^2 + \left(f(\mathbf{x})^T f(\mathbf{x}) - 2f(\mathbf{x})^T H^T f(\mathbf{x}) + f(\mathbf{x})^T H^T H f(\mathbf{x})\right) \\
 &= \left(n - \text{tr}[H] + \text{tr}[H^T H]\right) \sigma^2 + (f(\mathbf{x}) - Hf(\mathbf{x}))^T (f(\mathbf{x}) - Hf(\mathbf{x})),
 \end{aligned}$$

so the estimator

$$\hat{\sigma}^2 = \frac{\|r^2\|_2^2}{n - \text{tr}[H] + \text{tr}[H^T H]}$$

will be nearly unbiased in σ^2 when $f(\mathbf{x}) \approx Hf(\mathbf{x})$.

(E)

(F)

(G)

Problem 5

Gaussian processes

(A)

Problem 6

In nonparametric regression and spacial smoothing

(A)

R code for myfuns03.R

```
#####
##### Created by Spencer Woody on 11 Feb 2017 #####
#####

5 # =====
# Linear smoothing =====
# =====

lin.smooth <- function(x.new, x, y, kern.fun, h) {
10 # -----
# Linear smoother for some kernel function
# -----
# INPUTS:
# x.new – a single new point for which to estimate f(x.new)
15 # x – a vector of covariates from previous observations
# y – a vector of responses from previous observations
# kern.fun – some kernel function (e.g. Gaussian)
# *** takes 2 arguments: distance (dist) and bandwidth (h)
# h is the bandwidth for the kernel function
20 # -----
# OUTPUT:
# weights – a vector of weights for a new observation
# -----

25 weights <- kern.fun(dist = x - x.new, h = h) / h
weights <- weights / sum(weights)

fit <- crossprod(weights, y)

30 return(fit)
}

kern.unif <- function(dist, h) {
35 # -----
# Uniform kernel function
# -----
# INPUTS:
# dist
# h is
40 # Sigma is the covariance matrix
# -----
# OUTPUT:
# kern – the value of the uniform kernel function
# -----

45 kern <- ( abs(dist / h) <= 1) / 2

return(kern)
}

50 kern.norm <- function(dist, h) {
# -----
```

```

# Gaussian (normal) kernel function
# -----
55 # INPUTS:
# dist the distance
# h is the bandwidth
# Sigma is the covariance matrix
# -----
60 # OUTPUT:
# kern is the value of the Gaussian kernel function
# -----

kern <- 1 / sqrt(2 * pi) * exp(-(dist / h)^2 / 2)
65
return(kern)
}

make.noise <- function(x, f, res.dist, sd = NA, scale = NA, df = NA) {
70 # -----
# Simulate noisy response from some non-linear function
# -----
# INPUTS:
# x - the predictor values
75 # f - a function for the expected value,  $E(y | x) = f(x)$ 
# res.dist - a string for distribution of errors, either
#           "normal" for normal errors, or
#           "cauchy" for cauchy error
# sd - the standard deviation for residuals (for normal errors only)
80 # scale - the scale parameter for residuals (for cauchy errors only)
# -----
# OUTPUT:
# noise - the simulated noisy responses
# -----

85 if (res.dist == "normal" & !is.na(sd)) {
  noisy <- f(x) + rnorm(n = length(x), mean = 0, sd = sd)
} else if (res.dist == "cauchy" & !is.na(scale)) {
  noisy <- f(x) + rcauchy(n = length(x), location = 0, scale = scale)
90 } else {
  stop(paste("Must give res.dist argument as ",
             "either \"normal\" or \"cauchy\", AND ",
             "also specify sd (for normal) or ",
             "scale (for cauchy)", sep = ""))
95 }

return(noisy)
}

100 # =====
# Cross-validation =====
# =====

cv <- function(x.tr, y.tr, x.te, y.te, KERN.FUN, h) {
105 # -----

```

```

# Give cross validation mean squared prediction error
# -----
# INPUTS:
# x.tr – vector of predictors in * training * set
# y.tr – vector of responses in * training * set
# x.te – vector of predictors in * testing * set
# y.te – vector of responses in * testing * set
# KERN.FUN – some kernel function (e.g. Gaussian)
#           *** takes 2 arguments: distance (dist) and bandwidth (h)
# h – vector or bandwidths
# -----
# OUTPUT:
# mse – mean square predictive error
# -----

numbands <- length(h)
mse <- rep(NA, numbands)

for (i in 1:numbands) {
  y.pr <- sapply(
    x.te,
    lin.smooth,
    x = x.tr,
    y = y.tr,
    kern.fun = KERN.FUN,
    h = h[i]
  )

  mse[i] <- mean((y.te - y.pr)^2)
}

return(mse)
}

# =====
# Local polynomial regression =====
# =====

loc.lin <- function(x.new, x.vec, y.vec, h) {
  # -----
  # Give local linear estimator at some new point with normal kernel
  # -----
  # INPUTS:
  # x.new is some new point on x
  # x.vec – vector of x in sample
  # y.vec – vector of y in sample
  # h – bandwidth
  # -----
  # OUTPUT:
  # fit – the estimated value of f at x using local linear estimator
  # -----

  kern.x <- kern.norm(x.new - x.vec, h)

```

```

160   s1 <- sum(kern.x * (x - x.new))
      s2 <- sum(kern.x * (x - x.new)^2)
      w.x <- kern.x * (s2 * x.new - s1 * (x - x.new))

      fit <- crossprod(w.x, y.vec) / sum(w.x)
165
      return(fit)
    }

loc.pol <- function(x.new, x.vec, y.vec, D, h, give.mat = FALSE) {
170   # -----
      # Local polynomial reg. estimator at some new point with normal kernel
      # -----
      # INPUTS:
      # x.new is some new point on x
      # x.vec – vector of x in sample
175   # y.vec – vector of y in sample
      # D – dimension of the polynomial estimator (e.g., D = 1 means linear)
      # h – bandwidth for normal kernel
      # -----
180   # OUTPUT (list):
      # fit – the estimated value of f at x using local linear estimator
      # hatmat.vec – normalized weights, to be used in creating a hat matrix
      # -----

      # Number of observations
      N <- length(x.vec)

      # Vector of weights from Gaussian kernel
      w <- kern.norm(x.new - x.vec, h) / h
190

      # Create (non-normalized weights)
      if (D == 0) { # case of local constant estimator

        b <- matrix(w, nrow = 1)
195

      } else { # case of local polynomial estimator

        # Create R matrix
        R.x <- matrix(nrow = N, ncol = (D + 1))
        R.x[, 1] <- rep(1, N)
200

        for (j in 2:(D+1)) {
          R.x[, j] <- (x.new - x.vec)^(j - 1)
        }

        W.diag <- diag(w)

        RxTW <- crossprod(R.x, W.diag)

        b <- crossprod(w, R.x) %*% solve(RxTW %*% R.x) %*% RxTW
210
      }

```

```

fit <- tcrossprod((b / sum(b)), y.vec)

# Should the hat matrix vector be given?
if (give.mat) {
  output <- list("fit" = fit, "hatmat.vec" = b / sum(b))
} else {
  output <- fit
}

return(output)
}

loc.pol.hatmat <- function(x, y, D, h) {
  # -----
  # Create a hat matrix from local polynomial estimator
  # -----
  # INPUTS:
  # x - vector of x in sample
  # y - vector of y in sample
  # D - dimension of the polynomial estimator (e.g., D = 1 means linear)
  # h - bandwidth for normal kernel
  # -----
  # OUTPUT (list):
  # hatmat - hat matrix from local polynomial estimator
  # -----

  N <- length(x)

  hatmat <- matrix(nrow = N, ncol = N)

  for (i in 1:N) {
    hatmat[i, ] <- loc.pol(x[i], x, y, D, h, give.mat = TRUE)$hatmat.vec
  }

  return(hatmat)
}

loocv <- function(y, H) {
  # -----
  # Generic leave-one-out cross validation
  # -----
  # INPUTS:
  # y - the response vector
  # H - the "hat" matrix
  # -----
  # OUTPUT:
  # loocv - leave-one-out cross validation error
  # -----

  y.hat <- H %*% y

  loocv <- sum( ( {y - y.hat} / {1 - diag(H)} )^2 )

```

```

265     return(loocv)
}

# =====
270 # Gaussian process =====
# =====

my.mvn <- function(n, mu, Sigma) {
  # Simulate n draws from MVN(mu, Sigma)
275   #
  # Note: this function assumes that X already has an intercept term
  # (or doesn't, if we want to force OLS through the origin)
  #
  # INPUTS:
280   # n is the number of draws
  # mu is the mean vector
  # Sigma is the covariance matrix
  #
  # OUTPUT:
285   # x is matrix of n draws from MVN(mu, Sigma) [with n rows, p columns]
  #

  # dimension of MVN
  p <- length(mu)

290   # Check if inputs are valid (dimensions match, Sigma is square and p.s.d.)
  cond<- (ncol(Sigma) != p) |
          (nrow(Sigma) != p) |
          (max(eigen(Sigma)$values) <= 0)

295   if (cond) {
     return("Try again...")
   }

300   # Generate n*p univariate standard normal variables
  z <- matrix(rnorm(n*p), nrow = p)

  # Create a matrix containing copies of mu
  mumat <- matrix(rep(mu, n), nrow = p)

305   # Decompose Sigma into Sigma = L %*% Lt
  Lt <- chol(Sigma)

  # Generate sample with affine transformation of z
310   x <- crossprod(Lt, z) + mumat

  return(t(x))
}

315 ell2 <- function(x) {
  # Compute the ell2 norm of x, a vector in Euclidean space

```

```

    return(sqrt(sum(x^2)))
}

320 C.SE <- function(x.i, x.j, params = NA) {
    # -----
    # Compute the (i, j) element of a squared exp. covariance matrix
    # -----
325 # INPUTS:
    # x.i and x.j are two vectors in same space (need not be [0, 1])
    # params should be a vector of three hyperparameters
    # 1) b
330 # 2) tau1.sq
    # 3) tau2.sq
    # -----
    # OUTPUT:
335 # c.se is the value of the Matern-5/2 covariance matrix for x.i and x.j
    # -----

    if (prod(is.na(params))) {
        return("Must have three valid parameters.")
340 }

    if (length(params) != 3) {
        return("Must have three valid parameters.")
    }

345 b <- params[1]
    tau1.sq <- params[2]
    tau2.sq <- params[3]

350 b <- params[1]
    tau1.sq <- params[2]
    tau2.sq <- params[3]

355 # Euclidean distance between x.i and x.j
    d <- ell2(x.i - x.j)

    c.se <- tau1.sq * exp(-0.5 * (d / b)^2) + tau2.sq * (x.i == x.j)

360 return(c.se)
}

C.M52 <- function(x.i, x.j, params = NA) {
    # -----
365 # Compute the (i, j) element of a Matern-5/2 covariance matrix
    # -----
    # INPUTS:
    # x.i and x.j are two vectors in same space (need not be [0, 1])
370 # params should be a vector of three hyperparameters

```



```

#      1) b
#      2) tau1.sq
#      3) tau2.sq
#
375 # -----
# OUTPUT:
# c.m52 is the value of the Matern-5/2 covariance matrix for x.i and x.j
# -----

380 if (prod(is.na(params))) {
      return("Must have three valid parameters.")
    }

385 if (length(params) != 3) {
      return("Must have three valid parameters.")
    }

    b      <- params[1]
    tau1.sq <- params[2]
390 tau2.sq <- params[3]

    # Euclidean distance between x.i and x.j
    d <- ell2(x.i - x.j)

395 c.m52 <- tau1.sq * ( 1 + (5^0.5 * d / b) + (5 / 3 * (d / b)^2) ) *
      exp(-5^0.5 * d / b) + tau2.sq * (x.i == x.j)

    return(c.m52)
  }
400
make.covmat <- function(x, cov.fun, params = NA) {
  # -----
  # Compute the covariance matrix for a GP, given some cov. function
  #
405 # -----
  # INPUTS:
  # x is a vector of N values in [0, 1]
  # params should be a vector of three hyperparameters
  #      1) b
410 #      2) tau1.sq
  #      3) tau2.sq
  #
  # -----
  # OUTPUT:
415 # covmat is the covariance matrix of GP
  # -----

  if (prod(is.na(params))) {
    return("Must have three valid parameters.")
  }
420

  if (length(params) != 3) {
    return("Must have three valid parameters.")
  }

```

```
    }  
425  
    N <- length(x)  
  
    covmat <- matrix(nrow = N, ncol = N)  
  
430    for (j in 1:N) {  
      for (i in j:N) {  
        covmat[i, j] <- cov.fun(x[i], x[j], params = params)  
        covmat[j, i] <- covmat[i, j]  
      }  
435    }  
  
    return(covmat)  
  }
```

R code for exercises03.R

```
#####
##### Created by Spencer Woody on 11 Feb 2017 #####
#####

5 library(ggplot2)
  library(reshape2)
  library(gridExtra)
  library(wesanderson) # nice palettes

10 # Prep color palette
  # pal <- wes_palette("Zissou", 5)
  # col1 <- pal[5]
  # col2 <- pal[4]
  # col3 <- pal[1]

15 col1 <- "red"
  col2 <- "orange"
  col3 <- "blue"

20 source("myfuns03.R")

# =====
# Linear smoothers (part one) =====
# =====

25 # nonlinear function f(x)
  f1 <- function(x){
    return(x * (x - 4) * (x + 4))
  }

30 # Predictor vector
  x1 <- seq(-5, 5, length.out = 40)

  # Create sequence along x-space
35 x.seq <- seq(min(x1), max(x1), length.out = 200)

  # Response vector
  y1 <- make.noise(x1, f1, "normal", sd = 15)

40 # Bin width
  h1 <- 0.75

  # Gaussian kernel smoothing
  y.norm <- sapply(
45     x.seq,
     lin.smooth,
     x = x1,
     y = y1,
     kern.fun = kern.norm,
50     h = h1
  )
```

```
# Uniform kernel smoothing
y.unif <- sapply(
55   x.seq,
   lin.smooth,
   x = x1,
   y = y1,
   kern.fun = kern.unif,
60   h = h1
)

# Make a nice plot
h <- qplot(x.seq, geom = "blank") +
65 xlab("x") +
   ylab("y") +
   ggtitle(sprintf("Smoothing of cubic function")) +
   geom_point(aes(x = x1, y = y1), pch = 1) +
   stat_function(fun = f1, col = col1, linetype = "dashed") +
70 geom_line(aes(y = y.norm, colour = "Gaussian kernel")) +
   geom_line(aes(y = y.unif, colour = "Uniform kernel")) +
   scale_colour_manual(name = "Smoother", values = c(col3, col2)) +
   theme(legend.position = c(0.25, 0.15),
         text = element_text(family="Helvetica"))
75
h

pdf("firstexample.pdf")
h
80 dev.off()

# =====
85 # Linear smoothers (cross validation) ==== make big heat map =====
# =====

# Sample size
N = 500
90

# Set limits of x-space
xlo = 0
xhi = 1

95 # Number of bins for cross-validation
numbins <- 5

# Vector for bandwidths to search over
h.vec <- seq(0.001, 0.125, length.out = 100)
100

# Vector for standard deviations
s.vec <- seq(0.001, 0.5, length.out = 100)
s.vec <- rev(s.vec)

105 # Vector for periods
```

```

p.vec <- seq(0.1, 1, length.out = 100)

# Matrix for optimal bandwidth values
opt.h <- matrix(nrow = length(p.vec), ncol = length(s.vec))

110 for (i in 1:length(p.vec)) {
    # Select what the period is for this iteration
    p.i <- p.vec[i]

115     # Create a new function with this current period
    mysin.i <- function(x) {
        return(sin(x * 2*pi / p.i))
    }

120     for (j in 1:length(s.vec)) {
        # Select what the residual sd is for this iteration
        s.j <- s.vec[j]

        # Generate data
125         x.ij <- (xhi - xlo) * runif(N) + xlo
        y.ij <- make.noise(x.ij, mysin.i, "normal", sd = s.j)

        # Prepare mse matrix for this current iteration
        mse.ij <- matrix(nrow = numbins, ncol = length(h.vec))

130         # Create random partition into five bins
        jumble <- sample(1:N, N, replace = F)
        bin.indices <- split(jumble, cut(1:N, numbins))

135         for (bin in 1:numbins) {
            my.indices <- bin.indices[[bin]]
            x.tr <- x.ij[-(my.indices)]
            y.tr <- y.ij[-(my.indices)]
            x.te <- x.ij[my.indices]
            y.te <- y.ij[my.indices]
140             mse.ij[bin, ] <- cv(x.tr, y.tr, x.te, y.te, kern.norm, h.vec)
        }

        # Average out MSE over all bins
145         mse.vec <- colMeans(mse.ij)

        # Choose bandwidth with lowest average MSE
        opt.h[i, j] <- h.vec[which.min(mse.vec)]

150     }
    print(i)
}

# Check to make sure optimal bandwidths look OK
155 opt.h / max(opt.h)

# Plot these things
ohm <- melt(opt.h)

```

```

160 w <- ggplot(ohm, aes(rev(Var1), Var2, z = value)) +
  ggtitle("Picking Optimal Bandwidth for Gaussian Kernel Smoothing (5-fold CV)") +
  xlab("Decreasing standard deviation (0.5:0.001) (less noisy)") +
  ylab("Decreasing period (1:0.1) (more wiggly)") +
  geom_tile(aes(fill = value)) +
165 scale_fill_distiller("Bandwidth", palette = "Spectral")
  w

  # Save this to PDF
  pdf("img/opth.pdf", width = 7, height = 6)
170 w
  dev.off()

  # =====
175 # Linear smoothers (cross validation) ==== make 2 x 2 plot =====
  # =====

  # Sample size
  N = 500

180 # Set limits of x-space
  xlo = 0
  xhi = 1

185 # Number of bins for cross-validation
  numbins <- 5

  # Vector for bandwidths to search over
  h.vec <- seq(0.001, 0.125, length.out = 100)
190 # Vector for standard deviations
  # s.vec <- seq(0.001, 0.5, length.out = 64)
  s.vec <- c(0.1, 0.5)

195 # Vector for periods
  # p.vec <- seq(0.1, 1, length.out = 64)
  p.vec <- c(0.125, 1)

  # Matrix for optimal bandwidth values
200 opt.h <- matrix(nrow = length(p.vec), ncol = length(s.vec))

  # Matrix for random x-values and y-values
  x.mat <- matrix(nrow = nrow(opt.h) * ncol(opt.h), ncol = N)
  y.mat <- matrix(nrow = nrow(opt.h) * ncol(opt.h), ncol = N)
205 # Matrix for values of f at x
  x.seq <- seq(xlo, xhi, length.out = 200)
  fx.mat <- matrix(nrow = nrow(opt.h) * ncol(opt.h), ncol = length(x.seq))
  smooth <- matrix(nrow = nrow(opt.h) * ncol(opt.h), ncol = length(x.seq))
210 count <- 1

```

```

for (i in 1:length(p.vec)) {
  # Select what the period is for this iteration
215 p.i <- p.vec[i]

  # Create a new function with this current period
  mysin.i <- function(x) {
    return(sin(x * 2*pi / p.i))
220 }

  for (j in 1:length(s.vec)) {
    # Select what the residual sd is for this iteration
    s.j <- s.vec[j]
225

    x.ij <- (xhi - xlo) * runif(N) + xlo
    y.ij <- make.noise(x.ij, mysin.i, "normal", sd = s.j)

    fx.mat[count, ] <- mysin.i(x.seq)
230

    x.mat[count, ] <- x.ij
    y.mat[count, ] <- y.ij

    mse.ij <- matrix(nrow = numbins, ncol = length(h.vec))
235

    jumble <- sample(1:N, N, replace = F)
    bin.indices <- split(jumble, cut(1:N, numbins))

    for (bin in 1:numbins) {
      my.indices <- bin.indices[[bin]]
      x.tr <- x.ij[-(my.indices)]
      y.tr <- y.ij[-(my.indices)]
      x.te <- x.ij[my.indices]
      y.te <- y.ij[my.indices]
245 mse.ij[bin, ] <- cv(x.tr, y.tr, x.te, y.te, kern.norm, h.vec)
    }

    mse.vec <- colMeans(mse.ij)

    # Choose first 4/5 of x's and y's to be training, the rest of testing
    # x.tr <- x.ij[1:round(N*4/5)]
    # y.tr <- y.ij[1:round(N*4/5)]
    # x.te <- x.ij[-(1:round(N*4/5))]
    # y.te <- y.ij[-(1:round(N*4/5))]
255

    # mse.vec <- cv(x.tr, y.tr, x.te, y.te, kern.norm, h.vec)

    opt.h[i, j] <- h.vec[which.min(mse.vec)]

    smooth[count, ] <- y.norm <- sapply(
      x.seq,
      lin.smooth,
      x = x.ij,
      y = y.ij,
260

```

```

265         kern.fun = kern.norm,
           h = opt.h[i, j]
           )

270     count <- count + 1
       }
       print(i)
     }

275 y.min <- min(y.mat)
     y.max <- max(y.mat)

     ess <- qplot(x.seq, geom = "blank") +
       xlab("x") +
280       ylab("y") +
       ylim(y.min, y.max) +
       ggtitle(sprintf("sd = %5.3f, T = %5.3f, h = %5.5f", s.vec[1], p.vec[1], opt.h[1, 1])) +
       geom_point(aes(x = x.mat[1, ], y = y.mat[1, ]), pch = 1) +
       geom_line(aes(y = fx.mat[1, ]), col = "red", linetype = "dashed") +
285       geom_line(aes(y = smooth[1, ], colour = "Gaussian kernel")) +
       scale_colour_manual(name = "Smoother", values = "blue") +
       theme(legend.position = c(0.25, 0.15), text = element_text(family="Helvetica"))

     tee <- qplot(x.seq, geom = "blank") +
290     xlab("x") +
       ylab("y") +
       ylim(y.min, y.max) +
       ggtitle(sprintf("sd = %5.3f, T = %5.3f, h = %5.5f", s.vec[2], p.vec[1], opt.h[1, 2])) +
       geom_point(aes(x = x.mat[2, ], y = y.mat[2, ]), pch = 1) +
295       geom_line(aes(y = fx.mat[2, ]), col = "red", linetype = "dashed") +
       geom_line(aes(y = smooth[2, ], colour = "Gaussian kernel")) +
       scale_colour_manual(name = "Smoother", values = "blue") +
       theme(legend.position = c(0.25, 0.15), text = element_text(family="Helvetica"))

300     you <- qplot(x.seq, geom = "blank") +
       xlab("x") +
       ylab("y") +
       ylim(y.min, y.max) +
       ggtitle(sprintf("sd = %5.3f, T = %5.3f, h = %5.5f", s.vec[1], p.vec[2], opt.h[2, 1])) +
305       geom_point(aes(x = x.mat[3, ], y = y.mat[3, ]), pch = 1) +
       geom_line(aes(y = fx.mat[3, ]), col = "red", linetype = "dashed") +
       geom_line(aes(y = smooth[3, ], colour = "Gaussian kernel")) +
       scale_colour_manual(name = "Smoother", values = "blue") +
       theme(legend.position = c(0.25, 0.15), text = element_text(family="Helvetica"))

310     vee <- qplot(x.seq, geom = "blank") +
       xlab("x") +
       ylab("y") +
       ylim(y.min, y.max) +
315       ggtitle(sprintf("sd = %5.3f, T = %5.3f, h = %5.5f", s.vec[2], p.vec[2], opt.h[2, 2])) +
       geom_point(aes(x = x.mat[4, ], y = y.mat[4, ]), pch = 1) +
       geom_line(aes(y = fx.mat[4, ]), col = "red", linetype = "dashed") +

```



```

geom_line(aes(y = smooth[4, ], colour = "Gaussian kernel")) +
scale_colour_manual(name = "Smoother", values = "blue") +
320 theme(legend.position = c(0.25, 0.15), text = element_text(family="Helvetica"))

pdf("img/2x2.pdf")
grid.arrange(tee, ess, vee, you)
325 dev.off()

# =====
# Local polynomial regression =====
# =====
330 # -----
# Read in the data -----
# -----

335 utilities <- read.csv("utilities.csv", header = T)

x <- utilities$temp
y <- log(utilities$gasbill / utilities$billingdays)
# y <- log(utilities$gasbill)
340 # -----
# Cross validation -----
# -----

345 h.vec <- seq(0.25, 20, length.out = 500)
num.h <- length(h.vec)

hatmat.list <- list()
loocv.vec <- rep(NA, num.h)
350 for (i in 1:num.h) {
  hatmat.i <- loc.pol.hatmat(x, y, D = 1, h = h.vec[i])
  hatmat.list[[i]] <- hatmat.i

  355 loocv.vec[i] <- loocv(y, hatmat.i)

  if ((i %% 20) == 0) {
    print(i)
  }
360 }

plot(h.vec, loocv.vec, type = "l")

h.opt <- h.vec[which.min(loocv.vec)]
365 # -----
# Make a plot -----
# -----

370 x.seq <- seq(min(x), max(x), length.out = 200)

```

```

y.smooth <- sapply(
  x.seq,
  loc.pol,
  x.vec = x,
  y.vec = y,
  D = 1,
  h = h.opt
)

# Fitted y values
Hatmat <- loc.pol.hatmat(x, y, D = 1, h = h.opt)
y.hat <- Hatmat %*% y

var.est <- sum((y - y.hat)^2) /
(length(x) - sum(diag(Hatmat)) + sum(diag(crossprod(Hatmat))))

y.lo <- y.smooth - 1.96 * sqrt(var.est)
y.hi <- y.smooth + 1.96 * sqrt(var.est)

resplot <- qplot(x, y - y.hat) +
xlab(expression(paste("Temperature (",degree,"F)"))) +
ylab("residual") +
ggtitle("Residual plot for log-transformed data")

pdf("resplot2.pdf")
resplot
dev.off()

h <- qplot(x.seq, geom = "blank") +
xlab(expression(paste("Temperature (",degree,"F)"))) +
ylab("log-Daily gas bill (USD)") +
labs(title = "Daily gas bills for single-family homes in Minnesota") +
geom_ribbon(aes(ymin = y.lo, ymax = y.hi), fill = "grey80") +
geom_point(aes(x = x, y = y), pch = 1) +
geom_line(aes(y = y.smooth, colour = sprintf("h = %5.4f", h.opt))) +
scale_colour_manual(name = "Bandwidth", values = "firebrick3") +
theme(plot.title = element_text(hjust = 0.5),
text = element_text(family = "Trebuchet MS"),
legend.position = c(0.25, 0.15))

h

# R-squared
# 1 - sum((loc.pol.hatmat(x, y, 1, 5) %*% y - y)^2) / sum((y - mean(y))^2)

# =====
# Gaussian process =====
# =====

x.seq <- seq(0, 1, length.out = 100)

```

```

b <- 1
425 tau1.sq <- 1e-6
    tau2.sq <- 1e-5

myparams <- c(b, tau1.sq, tau2.sq)

430 xCM52 <- make.covmat(x.seq, C.M52, params = myparams)
    xSE <- make.covmat(x.seq, C.SE, params = myparams)

# =====
435 # Extra code for CV plotting.... =====
# =====

# Cross-validation with error bars

440 RSS.vec <- apply(RSS.mat, 2, mean)
    RSS.SE <- apply(RSS.mat, 2, sd) / sqrt(numbins)

lower = RSS.vec - RSS.SE
upper = RSS.vec + RSS.SE
445

# Make a plot!
pdf("perror.pdf", width = 12 / 1.25, height = 8 / 1.25)
h <- qplot(log(fit1$lambda), RSS.vec, geom = "path")
h + xlab(expression(paste("Penalization term, log(", lambda, ")"))) +
450 ylab("Expected prediction error") +
    labs(title = sprintf("%i-fold Cross-validation, Mallow's CP, and In-sample MSE", numbins)) +
    geom_ribbon(aes(ymin = lower, ymax = upper), fill = "grey80") +
    geom_line(aes(y = RSS.vec, colour = "CV"), show.legend = TRUE) +
    geom_line(aes(y = MSE.lasso, colour = "In-sample MSE"), show.legend = TRUE) +
455 geom_line(aes(y = MCP, colour = "CP"), show.legend = TRUE) +
    geom_vline(xintercept = log(minlambdaCV), linetype = 3, col = "black", size = 0.75, show.legend = TRUE)
    geom_vline(xintercept = log(minlambdaMCP), linetype = 3, col = "red", size = 0.75, show.legend = TRUE)
    scale_colour_manual(name = "", values = c("CV" = "black", "In-sample MSE" = "blue", "CP" = "red"))
    dev.off()
460

numbins <- 10
jumble <- sample(1:N2, N2, replace = F)
bin.indices <- split(jumble, cut(1:N2, numbins))

```