

# **SDS 383D: Exercises 3 – Linear smoothing and Gaussian processes**

February 22, 2017

*Professor Scott*

**Spencer Woody**

## **Problem 1**

### **Basic Concepts**

(A)

## Problem 2

### Curve fitting by linear smoothing

In this problem, consider a general nonlinear regression with one predictor and one response,  $y_i = f(x_i) + \epsilon_i$ , where  $\epsilon_i$  are mean-zero random variables.

- (A) For now, consider a linear regression on a response  $y_i$  with one predictor  $x_i$ , and both  $y_i$  and  $x_i$  have had their means subtracted, so the  $y_i = \beta x_i + \epsilon_i$ . Define  $S_x := \sum_{i=1}^n x_i^2$ . The least squares estimate for the coefficient, from Exercises 1, is

$$\begin{aligned}\hat{\beta} &= (X^T X)^{-1} X^T y \\ &= (x^T x)^{-1} x^T y \\ &= \frac{\sum_{i=1}^n x_i \cdot y_i}{\sum_{i=1}^n x_i^2} \\ &= \frac{\sum_{i=1}^n x_i \cdot y_i}{S_x} \\ &= \sum_{i=1}^n \frac{x_i}{S_x} \cdot y_i.\end{aligned}$$

So now our prediction  $y^*|x^*$  is,

$$\begin{aligned}\hat{y}^* &= \hat{f}(x^*) \\ &= \hat{\beta} x^* \\ &= \left( \sum_{i=1}^n \frac{x_i}{S_x} \cdot y_i \right) \cdot x^* \\ &= \sum_{i=1}^n \left( \frac{x_i}{S_x} \cdot x^* \right) \cdot y_i,\end{aligned}$$

which we recognize as being in the form of the general *linear smoother*

$$\hat{f}(x^*) = \sum_{i=1}^n w(x_i, x^*) \cdot y_i$$

for some weight function  $w(x_i, x^*)$ . In particular, the weight function for linear regression gives weight to each  $y_i$  proportional to the value of  $x_i$ . Contrast this with the  $k$ -nearest neighbors smoothing weight function,

$$w_K(x_i, x^*) = \begin{cases} 1/K & \text{if } x_i \text{ is one of the } K \text{ closest sample points to } x^*, \\ 0 & \text{otherwise} \end{cases},$$

which gives *equal* weight to  $y_i$ s but *only* to the  $k$ -nearest neighbors of  $x^*$ .

- (B) Now we have the very general weight function

$$w(x_i, x^*) = \frac{1}{h} \cdot K\left(\frac{x_i - x^*}{h}\right)$$

where  $K(\bullet)$  is some kernel function. The script `myfuns03.R` in the appendix shows an R function for linear smoothing, as well functions for the uniform and Gaussian kernels. Figure 1 shows an example of smoothing with a bandwidth of 0.75 for a cubic function  $f(x)$  with iid residuals from the  $\mathcal{N}(0, 15^2)$  distribution.

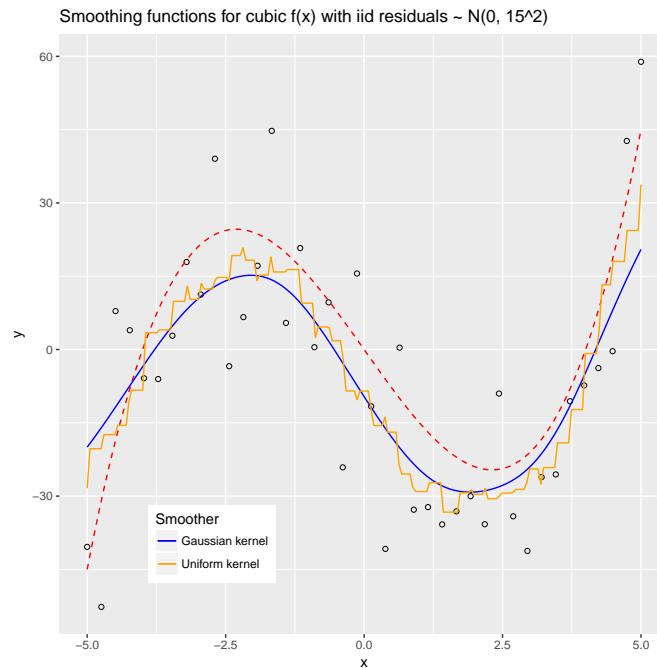


Figure 1: Uniform and Gaussian kernel smoothing for  $y = f(x) + e$ ,  $f(x) = x(x - 4)(x + 4)$ ,  $h = 0.75$

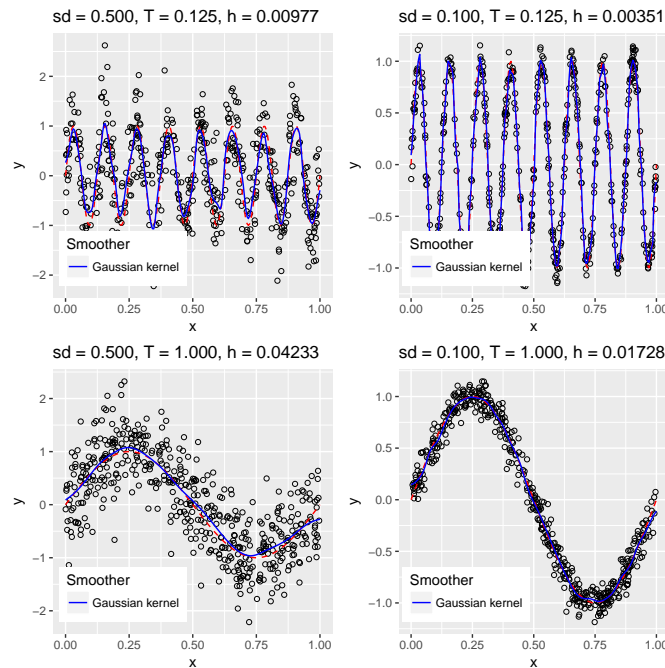
## Problem 3

### Cross validation

- (A) See attached R code for a script to return prediction error estimates for smoothing given a specified choice of bandwidth,  $h$ .
- (B) For this exercise, I produced 500 data points on the  $x$ -space  $[0, 1]$  from a sine function  $f(x)$  with a given period and set the amplitude, and added Gaussian noise with a given standard deviation. Then I used 5-fold cross validation to select the optimal bandwidth for that given period and standard deviation of noise term. Figure 2 shows the optimal bandwidths for period ranging from 0.1 to 1, and standard deviation ranging from 0.001 to 0.5, and Figure 3 shows four example. The highest bandwidths are chosen for functions with high “wigglyness” and high noise, and the smallest bandwidths are chosen for functions with low “wigglyness” and low noise. This makes sense. As the frequency increases (i.e., period decreases) then we need a tighter bandwidth because the value of the function is fluctuating at a greater rate. As noise increases, we need a greater bandwidth to smooth out the noise. Furthermore, we can see that in all cases we recover the underlying function pretty well.

Figure 2: Optimal bandwidths for varying periods and standard deviations

- (C) I’ll get around to this eventually. . . .


 Figure 3:  $2 \times 2$  example with fitted curves

## Problem 4

### Local polynomial regression

(A) Define

$$g_x(u; a) = a_0 + \sum_{k=1}^D a_k (u - x)^k \quad (1)$$

$$= \begin{cases} \sum_{j=0}^{D+1} a_j (u - x)^j & \text{if } u \neq x \\ a_0 & \text{if } u = x \end{cases} \quad (2)$$

The coefficients of  $a$  will come from the weighted least squares problem

$$\hat{a} = \arg \min_{a \in \mathcal{R}^{D+1}} \sum_{i=1}^n w_i [y_i - g_x(x_i, a)]^2 \quad (3)$$

Furthermore, define the matrix  $R_x$  whose  $(i, j)$  element is  $(x_i - x)^2$ . Then the estimate  $\hat{f}(x)$  will be  $R_x \hat{a}$ . The solution of  $\hat{a}$  may be found with

$$\begin{aligned} \hat{a} &= \arg \min_{a \in \mathcal{R}^{D+1}} (y - R_x a)^T W (y - R_x a) \\ &= (R_x^T W R_x)^{-1} R_x^T W y \end{aligned}$$

where  $W = \text{diag}(w_1, \dots, w_n)$  following the same argument to find the WLS estimate of linear regression from Exercises 1.

(B)

(C)

- (D) With  $H$  a smoothing matrix (or “hat matrix”), let  $r = y - Hy$  be the vector of residuals. If the random vector  $x$  with mean vector  $\mu$  and covariance matrix  $\Sigma$ , then  $E(x^T Q x) = \text{tr}(Q \Sigma) + \mu^T Q \mu$ . Then,

$$E \left( \|r^2\| \right) = E \left( (y - Hy)^T (y - Hy) \right)$$

(E)

(F)

(G)

## **Problem 5**

### **Gaussian processes**

(A)

## **Problem 6**

**In nonparametric regression and spacial smoothing**

(A)



R code for myfuns03.R

```
#####
##### Created by Spencer Woody on 11 Feb 2017 #####
#####

5 # =====
# Linear smoothing =====
# =====

lin.smooth <- function(x.new, x, y, kern.fun, h) {
10 # -----
# Linear smoother for some kernel function
# -----
# INPUTS:
# x.new – a new point for which to estimate f(x.new)
15 # x – a vector of covariates from previous observations
# y – a vector of responses from previous observations
# kern.fun – some kernel function (e.g. Gaussian)
# *** takes 2 arguments: distance (dist) and bandwidth (h)
# h is the bandwidth for the kernel function
20 # -----
# OUTPUT:
# weights – a vector of weights for a new observation
# -----

25 weights <- kern.fun(dist = x - x.new, h = h) / h
weights <- weights / sum(weights)

fit <- crossprod(weights, y)

30 return(fit)
}

kern.unif <- function(dist, h) {
35 # -----
# Uniform kernel function
# -----
# INPUTS:
# dist
# h is
40 # Sigma is the covariance matrix
# -----
# OUTPUT:
# kern – the value of the uniform kernel function
# -----

45 kern <- ( abs(dist / h) <= 1) / 2

return(kern)
}

50 kern.norm <- function(dist, h) {
# -----
```

```

# Gaussian (normal) kernel function
# -----
55 # INPUTS:
# dist the distance
# h is the bandwidth
# Sigma is the covariance matrix
# -----
60 # OUTPUT:
# kern is the value of the Gaussian kernel function
# -----

kern <- 1 / sqrt(2 * pi) * exp(-(dist / h)^2 / 2)
65
return(kern)
}

make.noise <- function(x, f, res.dist, sd = NA, scale = NA, df = NA) {
70 # -----
# Simulate noisy response from some non-linear function
# -----
# INPUTS:
# x - the predictor values
75 # f - a function for the expected value, E(y | x) = f(x)
# res.dist - a string for distribution of errors, either
#           "normal" for normal errors, or
#           "cauchy" for cauchy error
# sd - the standard deviation for residuals (for normal errors only)
80 # scale - the scale parameter for residuals (for cauchy errors only)
# -----
# OUTPUT:
# noise - the simulated noisy responses
# -----
85
if (res.dist == "normal" & !is.na(sd)) {
  noisy <- f(x) + rnorm(n = length(x), mean = 0, sd = sd)
} else if (res.dist == "cauchy" & !is.na(scale)) {
  noisy <- f(x) + rcauchy(n = length(x), location = 0, scale = scale)
90 } else {
  stop(paste("Must give res.dist argument as ",
             "either \"normal\" or \"cauchy\", AND ",
             "also specify sd (for normal) or ",
             "scale (for cauchy)", sep = ""))
95 }

return(noisy)
}

# =====
# Cross-validation =====
# =====

cv <- function(x.tr, y.tr, x.te, y.te, KERN.FUN, h) {
105 # -----

```

```

# Give cross validation mean squared prediction error
# -----
# INPUTS:
# x.tr – vector of predictors in * training * set
# y.tr – vector of responses in * training * set
# x.te – vector of predictors in * testing * set
# y.te – vector of responses in * testing * set
# kern.fun – some kernel function (e.g. Gaussian)
#           *** takes 2 arguments: distance (dist) and bandwidth (h)
# h – vector or bandwidths
# -----
# OUTPUT:
# mse – mean square predictive error
# -----

numbands <- length(h)
mse <- rep(NA, numbands)

for (i in 1:numbands) {
  y.pr <- sapply(
    x.te,
    lin.smooth,
    x = x.tr,
    y = y.tr,
    kern.fun = KERN.FUN,
    h = h[i]
  )

  mse[i] <- mean((y.te - y.pr)^2)
}

return(mse)
}

# =====
# Local polynomial regression =====
# =====

loc.lin <- function() {
  # -----
  # Give local linear estimator at some new point
  # -----
  # INPUTS:
  # x is some new point on x
  # x.vec – vector of x in sample
  # y.vec – vector of y in sample
  # kern.fun – some kernel function (e.g. Gaussian)
  #           *** takes 2 arguments: distance (dist) and bandwidth (h)
  # h – vector or bandwidths
  # -----
  # OUTPUT:
  # mse – mean square predictive error
  # -----

```

```

160     kern.x <- kern.norm(x - x.vec, h)

    w.x <-
    s1 <- sum(kern.x * (x.vec - x))

165 }

# =====
# Gaussian process =====
# =====

170 my.mvn <- function(n, mu, Sigma) {
    # Simulate n draws from MVN(mu, Sigma)
    #
    # Note: this function assumes that X already has an intercept term
175 # (or doesn't, if we want to force OLS through the origin)
    #
    # INPUTS:
    # n is the number of draws
    # mu is the mean vector
180 # Sigma is the covariance matrix
    #
    # OUTPUT:
    # x is matrix of n draws from MVN(mu, Sigma) [with n rows, p columns]
    #

185 # dimension of MVN
    p <- length(mu)

    # Check if inputs are valid (dimensions match, Sigma is square and p.s.d.)
190 cond<- (ncol(Sigma) != p) |
           (nrow(Sigma) != p) |
           (max(eigen(Sigma)$values) <= 0)

    if (cond) {
195         return("Try again...")
    }

    # Generate n*p univariate standard normal variables
    z <- matrix(rnorm(n*p), nrow = p)

200 # Create a matrix containing copies of mu
    mumat <- matrix(rep(mu, n), nrow = p)

    # Decompose Sigma into Sigma = L %*% Lt
205 Lt <- chol(Sigma)

    # Generate sample with affine transformation of z
    x <- crossprod(Lt, z) + mumat

210     return(t(x))
}

```

```

ell2 <- function(x) {
  # Compute the ell2 norm of x, a vector in Euclidean space
  return(sqrt(sum(x^2)))
}

C.SE <- function(x.i, x.j, params = NA) {
  # -----
  # Compute the (i, j) element of a squared exp. covariance matrix
  # -----
  # INPUTS:
  # x.i and x.j are two vectors in same space (need not be [0, 1])
  # params should be a vector of three hyperparameters
  #   1) b
  #   2) tau1.sq
  #   3) tau2.sq
  # -----
  # OUTPUT:
  # c.se is the value of the Matern-5/2 covariance matrix for x.i and x.j
  # -----

  if (prod(is.na(params))) {
    return("Must have three valid parameters.")
  }

  if (length(params) != 3) {
    return("Must have three valid parameters.")
  }

  b <- params[1]
  tau1.sq <- params[2]
  tau2.sq <- params[3]

  b <- params[1]
  tau1.sq <- params[2]
  tau2.sq <- params[3]

  # Euclidean distance between x.i and x.j
  d <- ell2(x.i - x.j)

  c.se <- tau1.sq * exp(-0.5 * (d / b)^2) + tau2.sq * (x.i == x.j)

  return(c.se)
}

C.M52 <- function(x.i, x.j, params = NA) {
  # -----
  # Compute the (i, j) element of a Matern-5/2 covariance matrix
  # -----

```

```

265 # -----
# INPUTS:
# x.i and x.j are two vectors in same space (need not be [0, 1])
# params should be a vector of three hyperparameters
#     1) b
270 #     2) tau1.sq
#     3) tau2.sq
#
# -----
# OUTPUT:
275 # c.m52 is the value of the Matern-5/2 covariance matrix for x.i and x.j
# -----

if (prod(is.na(params))) {
  return("Must have three valid parameters.")
280 }

if (length(params) != 3) {
  return("Must have three valid parameters.")
}

285 b      <- params[1]
tau1.sq <- params[2]
tau2.sq <- params[3]

290 # Euclidean distance between x.i and x.j
d <- ell2(x.i - x.j)

c.m52 <- tau1.sq * ( 1 + (5^0.5 * d / b) + (5 / 3 * (d / b)^2) ) *
  exp(-5^0.5 * d / b) + tau2.sq * (x.i == x.j)
295

return(c.m52)
}

make.covmat <- function(x, cov.fun, params = NA) {
300 # -----
# Compute the covariance matrix for a GP, given some cov. function
#
# -----
# INPUTS:
305 # x is a vector of N values in [0, 1]
# params should be a vector of three hyperparameters
#     1) b
#     2) tau1.sq
#     3) tau2.sq
310 #
# -----
# OUTPUT:
# covmat is the covariance matrix of GP
# -----

315 if (prod(is.na(params))) {
  return("Must have three valid parameters.")
}

```

```
    }  
320   if (length(params) != 3) {  
       return("Must have three valid parameters.")  
    }  
  
    N <- length(x)  
325  
    covmat <- matrix(nrow = N, ncol = N)  
  
    for (j in 1:N) {  
      for (i in j:N) {  
330        covmat[i, j] <- cov.fun(x[i], x[j], params = params)  
        covmat[j, i] <- covmat[i, j]  
      }  
    }  
  
335   return(covmat)  
}
```

## R code for exercises03.R

```
#####
##### Created by Spencer Woody on 11 Feb 2017 #####
#####

5 library(ggplot2)
  library(reshape2)
  library(gridExtra)
  library(wesanderson) # nice palettes

10 # Prep color palette
  # pal <- wes_palette("Zissou", 5)
  # col1 <- pal[5]
  # col2 <- pal[4]
  # col3 <- pal[1]

15 col1 <- "red"
  col2 <- "orange"
  col3 <- "blue"

20 source("myfuns03.R")

# =====
# Linear smoothers (part one) =====
# =====

25 # nonlinear function f(x)
  f1 <- function(x){
    return(x * (x - 4) * (x + 4))
  }

30 # Predictor vector
  x1 <- seq(-5, 5, length.out = 40)

  # Create sequence along x-space
35 x.seq <- seq(min(x1), max(x1), length.out = 200)

  # Response vector
  y1 <- make.noise(x1, f1, "normal", sd = 15)

40 # Bin width
  h1 <- 0.75

  # Gaussian kernel smoothing
  y.norm <- sapply(
45     x.seq,
     lin.smooth,
     x = x1,
     y = y1,
     kern.fun = kern.norm,
50     h = h1
  )
```



```

# Uniform kernel smoothing
y.unif <- sapply(
55   x.seq,
    lin.smooth,
    x = x1,
    y = y1,
    kern.fun = kern.unif,
60   h = h1
)

# Make a nice plot
h <- qplot(x.seq, geom = "blank") +
65 xlab("x") +
  ylab("y") +
  ggtitle(sprintf("Smoothing of cubic function")) +
  geom_point(aes(x = x1, y = y1), pch = 1) +
  stat_function(fun = f1, col = col1, linetype = "dashed") +
70 geom_line(aes(y = y.norm, colour = "Gaussian kernel")) +
  geom_line(aes(y = y.unif, colour = "Uniform kernel")) +
  scale_colour_manual(name = "Smoother", values = c(col3, col2)) +
  theme(legend.position = c(0.25, 0.15),
        text = element_text(family="Helvetica"))
75
h

pdf("firstexample.pdf")
h
80 dev.off()

# =====
# Linear smoothers (cross validation) ==== make big heat map =====
# =====
85

# Sample size
N = 500

# Set limits of x-space
90 xlo = 0
  xhi = 1

# Number of bins for cross-validation
numbins <- 5
95

# Vector for bandwidths to search over
h.vec <- seq(0.001, 0.125, length.out = 100)

# Vector for standard deviations
100 s.vec <- seq(0.001, 0.5, length.out = 100)
    s.vec <- rev(s.vec)

# Vector for periods
p.vec <- seq(0.1, 1, length.out = 100)
105

```

```

# Matrix for optimal bandwidth values
opt.h <- matrix(nrow = length(p.vec), ncol = length(s.vec))

for (i in 1:length(p.vec)) {
  # Select what the period is for this iteration
  p.i <- p.vec[i]

  # Create a new function with this current period
  mysin.i <- function(x) {
    return(sin(x * 2*pi / p.i))
  }

  for (j in 1:length(s.vec)) {
    # Select what the residual sd is for this iteration
    s.j <- s.vec[j]

    # Generate data
    x.ij <- (xhi - xlo) * runif(N) + xlo
    y.ij <- make.noise(x.ij, mysin.i, "normal", sd = s.j)

    # Prepare mse matrix for this current iteration
    mse.ij <- matrix(nrow = numbins, ncol = length(h.vec))

    # Create random partition into five bins
    jumble <- sample(1:N, N, replace = F)
    bin.indices <- split(jumble, cut(1:N, numbins))

    for (bin in 1:numbins) {
      my.indices <- bin.indices[[bin]]
      x.tr <- x.ij[-(my.indices)]
      y.tr <- y.ij[-(my.indices)]
      x.te <- x.ij[my.indices]
      y.te <- y.ij[my.indices]
      mse.ij[bin, ] <- cv(x.tr, y.tr, x.te, y.te, kern.norm, h.vec)
    }

    # Average out MSE over all bins
    mse.vec <- colMeans(mse.ij)

    # Choose bandwidth with lowest average MSE
    opt.h[i, j] <- h.vec[which.min(mse.vec)]
  }
  print(i)
}

# Check to make sure optimal bandwidths look OK
opt.h / max(opt.h)

# Plot these things
ohm <- melt(opt.h)

w <- ggplot(ohm, aes(rev(Var1), Var2, z = value)) +

```

```
ggtitle("Picking Optimal Bandwidth for Gaussian Kernel Smoothing (5-fold CV)") +
160 xlab("Decreasing standard deviation (0.5:0.001) (less noisy)") +
ylab("Decreasing period (1:0.1) (more wiggly)") +
geom_tile(aes(fill = value)) +
scale_fill_distiller("Bandwidth", palette = "Spectral")
w
165
# Save this to PDF
pdf("oph3.pdf", width = 7, height = 6)
w
dev.off()
170

# =====
# Linear smoothers (cross validation) ==== make 2 x 2 plot =====
175 # =====

# Sample size
N = 500

180 # Set limits of x-space
xlo = 0
xhi = 1

# Number of bins for cross-validation
185 numbins <- 5

# Vector for bandwidths to search over
h.vec <- seq(0.001, 0.125, length.out = 100)

190 # Vector for standard deviations
# s.vec <- seq(0.001, 0.5, length.out = 64)
s.vec <- c(0.1, 0.5)

# Vector for periods
195 # p.vec <- seq(0.1, 1, length.out = 64)
p.vec <- c(0.125, 1)

# Matrix for optimal bandwidth values
opt.h <- matrix(nrow = length(p.vec), ncol = length(s.vec))
200

# Matrix for random x-values and y-values
x.mat <- matrix(nrow = nrow(opt.h) * ncol(opt.h), ncol = N)
y.mat <- matrix(nrow = nrow(opt.h) * ncol(opt.h), ncol = N)

205 # Matrix for values of f at x
x.seq <- seq(xlo, xhi, length.out = 200)
fx.mat <- matrix(nrow = nrow(opt.h) * ncol(opt.h), ncol = length(x.seq))
smooth <- matrix(nrow = nrow(opt.h) * ncol(opt.h), ncol = length(x.seq))

210 count <- 1
```

```

for (i in 1:length(p.vec)) {
  # Select what the period is for this iteration
  p.i <- p.vec[i]
215

  # Create a new function with this current period
  mysin.i <- function(x) {
    return(sin(x * 2*pi / p.i))
  }
220

  for (j in 1:length(s.vec)) {
    # Select what the residual sd is for this iteration
    s.j <- s.vec[j]

225

    x.ij <- (xhi - xlo) * runif(N) + xlo
    y.ij <- make.noise(x.ij, mysin.i, "normal", sd = s.j)

    fx.mat[count, ] <- mysin.i(x.seq)

230

    x.mat[count, ] <- x.ij
    y.mat[count, ] <- y.ij

    mse.ij <- matrix(nrow = numbins, ncol = length(h.vec))

235

    jumble <- sample(1:N, N, replace = F)
    bin.indices <- split(jumble, cut(1:N, numbins))

    for (bin in 1:numbins) {
      my.indices <- bin.indices[[bin]]
240
      x.tr <- x.ij[-(my.indices)]
      y.tr <- y.ij[-(my.indices)]
      x.te <- x.ij[my.indices]
      y.te <- y.ij[my.indices]
      mse.ij[bin, ] <- cv(x.tr, y.tr, x.te, y.te, kern.norm, h.vec)
245
    }

    mse.vec <- colMeans(mse.ij)

    # Choose first 4/5 of x's and y's to be training, the rest of testing
250
    # x.tr <- x.ij[1:round(N*4/5)]
    # y.tr <- y.ij[1:round(N*4/5)]
    # x.te <- x.ij[-(1:round(N*4/5))]
    # y.te <- y.ij[-(1:round(N*4/5))]

255

    # mse.vec <- cv(x.tr, y.tr, x.te, y.te, kern.norm, h.vec)

    opt.h[i, j] <- h.vec[which.min(mse.vec)]

    smooth[count, ] <- y.norm <- sapply(
260
      x.seq,
      lin.smooth,
      x = x.ij,
      y = y.ij,
      kern.fun = kern.norm,

```

```

265         h = opt.h[i, j]
           )

           count <- count + 1
270     }
       print(i)
}

ess <- qplot(x.seq, geom = "blank") +
275 xlab("x") +
ylab("y") +
ggtitle(sprintf("sd = %5.3f, T = %5.3f, h = %5.5f", s.vec[1], p.vec[1], opt.h[1, 1])) +
geom_point(aes(x = x.mat[1, ], y = y.mat[1, ]), pch = 1) +
geom_line(aes(y = fx.mat[1, ]), col = "red", linetype = "dashed") +
280 geom_line(aes(y = smooth[1, ], colour = "Gaussian kernel")) +
scale_colour_manual(name = "Smoother", values = "blue") +
theme(legend.position = c(0.25, 0.15), text = element_text(family="Helvetica"))

tee <- qplot(x.seq, geom = "blank") +
285 xlab("x") +
ylab("y") +
ggtitle(sprintf("sd = %5.3f, T = %5.3f, h = %5.5f", s.vec[2], p.vec[1], opt.h[1, 2])) +
geom_point(aes(x = x.mat[2, ], y = y.mat[2, ]), pch = 1) +
geom_line(aes(y = fx.mat[2, ]), col = "red", linetype = "dashed") +
290 geom_line(aes(y = smooth[2, ], colour = "Gaussian kernel")) +
scale_colour_manual(name = "Smoother", values = "blue") +
theme(legend.position = c(0.25, 0.15), text = element_text(family="Helvetica"))

you <- qplot(x.seq, geom = "blank") +
295 xlab("x") +
ylab("y") +
ggtitle(sprintf("sd = %5.3f, T = %5.3f, h = %5.5f", s.vec[1], p.vec[2], opt.h[2, 1])) +
geom_point(aes(x = x.mat[3, ], y = y.mat[3, ]), pch = 1) +
geom_line(aes(y = fx.mat[3, ]), col = "red", linetype = "dashed") +
300 geom_line(aes(y = smooth[3, ], colour = "Gaussian kernel")) +
scale_colour_manual(name = "Smoother", values = "blue") +
theme(legend.position = c(0.25, 0.15), text = element_text(family="Helvetica"))

vee <- qplot(x.seq, geom = "blank") +
305 xlab("x") +
ylab("y") +
ggtitle(sprintf("sd = %5.3f, T = %5.3f, h = %5.5f", s.vec[2], p.vec[2], opt.h[2, 2])) +
geom_point(aes(x = x.mat[4, ], y = y.mat[4, ]), pch = 1) +
geom_line(aes(y = fx.mat[4, ]), col = "red", linetype = "dashed") +
310 geom_line(aes(y = smooth[4, ], colour = "Gaussian kernel")) +
scale_colour_manual(name = "Smoother", values = "blue") +
theme(legend.position = c(0.25, 0.15), text = element_text(family="Helvetica"))

315 pdf("2x2.pdf")
grid.arrange(tee, ess, vee, you)
dev.off()

```

```

# =====
320 # Gaussian process =====
# =====

x.seq <- seq(0, 1, length.out = 100)

325 b <- 1
tau1.sq <- 1e-6
tau2.sq <- 1e-5

myparams <- c(b, tau1.sq, tau2.sq)

330 xCM52 <- make.covmat(x.seq, C.M52, params = myparams)
xSE <- make.covmat(x.seq, C.SE, params = myparams)

335 # =====
# Extra code for CV plotting.... =====
# =====

# Cross-validation with error bars

340 RSS.vec <- apply(RSS.mat, 2, mean)
RSS.SE <- apply(RSS.mat, 2, sd) / sqrt(numbins)

lower = RSS.vec - RSS.SE
345 upper = RSS.vec + RSS.SE

# Make a plot!
pdf("perror.pdf", width = 12 / 1.25, height = 8 / 1.25)
h <- qplot(log(fit1$lambda), RSS.vec, geom = "path")
350 h + xlab(expression(paste("Penalization term, log(", lambda, ")"))) +
ylab("Expected prediction error") +
labs(title = sprintf("%i-fold Cross-validation, Mallow's CP, and In-sample MSE", numbins)) +
geom_ribbon(aes(ymin = lower, ymax = upper), fill = "grey80") +
geom_line(aes(y = RSS.vec, colour = "CV"), show.legend = TRUE) +
355 geom_line(aes(y = MSE.lasso, colour = "In-sample MSE"), show.legend = TRUE) +
geom_line(aes(y = MCP, colour = "CP"), show.legend = TRUE) +
geom_vline(xintercept = log(minlambdaCV), linetype = 3, col = "black", size = 0.75, show.legend = TRUE)
geom_vline(xintercept = log(minlambdaMCP), linetype = 3, col = "red", size = 0.75, show.legend = TRUE)
scale_colour_manual(name = "", values = c("CV" = "black", "In-sample MSE" = "blue", "CP" = "red"))
360 dev.off()

numbins <- 10
jumble <- sample(1:N2, N2, replace = F)
365 bin.indices <- split(jumble, cut(1:N2, numbins))

```